

SQL – SQL Joins

Dr. Villanes

Questions

Q1 – Q6

Question 1

With SQL, how do you select all the columns from a table named "Persons"?

A SELECT [all] FROM Persons

B SELECT Persons

C SELECT * FROM Persons

D SELECT *.Persons

Question 2

In the following query, what is the number 2 referring to?

```
select ID, Capital_Gain, capital_gain*0.10  
from exercise.records  
where capital_gain <> 0  
order by 2;
```

A ID

B Capital_Gain

C Capital_Gain*0.10

Question 3

With SQL, how can you return all the records from a table named "Persons" sorted descending by "FirstName"?

- A** `SELECT * FROM Persons ORDER BY FirstName DESC`
- B** `SELECT * FROM Persons ORDER FirstName DESC`
- C** `SELECT * FROM Persons SORT 'FirstName' DESC`
- D** `SELECT * FROM Persons SORT BY 'FirstName' DESC`

Question 4

With SQL, how can you return the number of records in the "Persons" table?

A SELECT COLUMNS(*) FROM Persons

B SELECT NO(*) FROM Persons

C SELECT COUNT(*) FROM Persons

D SELECT LEN(*) FROM Persons

Question 5

If the following query is submitted in the CUSTOMERS table, how many rows would be in the output?

```
SELECT NAME, SUM(SALARY)
FROM CUSTOMERS
GROUP BY NAME;
```

| ID | NAME | AGE | ADDRESS | SALARY |
|----|----------|-----|-----------|----------|
| 1 | Ramesh | 32 | Ahmedabad | 2000.00 |
| 2 | Khilan | 25 | Delhi | 1500.00 |
| 3 | kaushik | 23 | Kota | 2000.00 |
| 4 | Chaitali | 25 | Mumbai | 6500.00 |
| 5 | Hardik | 27 | Bhopal | 8500.00 |
| 6 | Komal | 22 | MP | 4500.00 |
| 7 | Muffy | 24 | Indore | 10000.00 |

A 0 rows

B 6 rows

C 7 rows

D 8 rows

Question 6

What about now?

```
SELECT NAME, SUM(SALARY)
FROM CUSTOMERS
GROUP BY NAME;
```

| ID | NAME | AGE | ADDRESS | SALARY |
|----|---------|-----|-----------|----------|
| 1 | Ramesh | 32 | Ahmedabad | 2000.00 |
| 2 | Ramesh | 25 | Delhi | 1500.00 |
| 3 | kaushik | 23 | Kota | 2000.00 |
| 4 | kaushik | 25 | Mumbai | 6500.00 |
| 5 | Hardik | 27 | Bhopal | 8500.00 |
| 6 | Komal | 22 | MP | 4500.00 |
| 7 | Muffy | 24 | Indore | 10000.00 |

A 0 rows

B 4 rows

C 5 rows

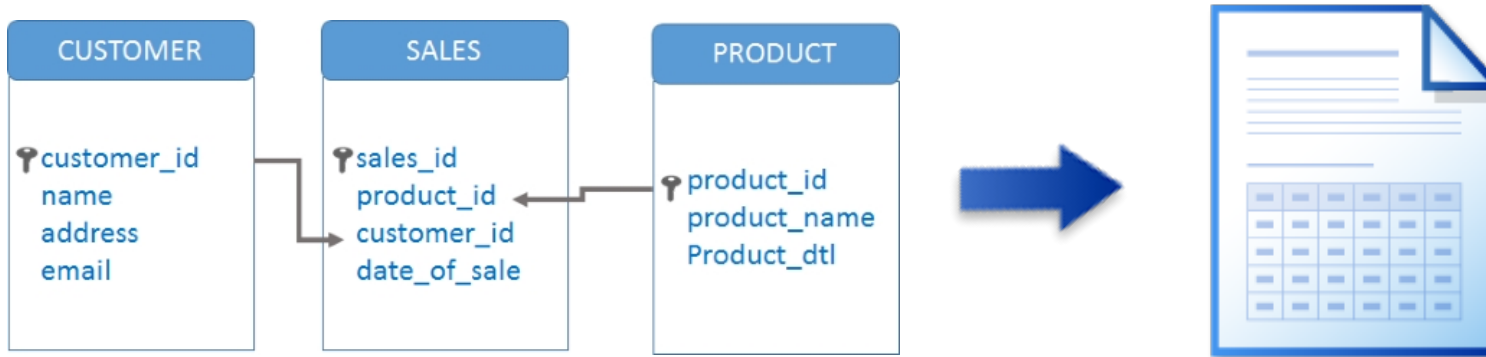
D 7 rows

Adding the new database Practice to DataGrip

SQL Joins


Combining Tables

SQL uses *joins* to combine tables horizontally. Requesting a join involves matching data from one row in one table with a corresponding row in a second table. Matching is typically performed on one or more columns in the two tables.




Types of Joins: two types


- *Inner joins* return only matching rows.




- *Outer joins* return all matching rows, plus nonmatching rows from one or both tables.



Left



Full



Right

Cartesian Product

A query that lists multiple tables in the FROM clause without a WHERE clause **produces all possible combinations of rows from all tables**. This result is called a *Cartesian product*.

```
select *  
from customers, transactions;
```


Cartesian Product

customers

| ID | Name |
|-----|-------|
| 101 | Smith |
| 104 | Jones |
| 102 | Blank |

transactions

| ID | Action | Amount |
|-----|----------|--------|
| 102 | Purchase | \$100 |
| 103 | Return | \$52 |
| 105 | Return | \$212 |



Result Set

**Non-
matching
IDs**

| ID | Name | ID | Action | Amount |
|-----|-------|-----|----------|--------|
| 101 | Smith | 102 | Purchase | \$100 |
| 101 | Smith | 103 | Return | \$52 |
| 101 | Smith | 105 | Return | \$212 |
| 104 | Jones | 102 | Purchase | \$100 |
| 104 | Jones | 103 | Return | \$52 |
| 104 | Jones | 105 | Return | \$212 |
| 102 | Blank | 102 | Purchase | \$100 |
| 102 | Blank | 103 | Return | \$52 |
| 102 | Blank | 105 | Return | \$212 |


9 rows

The Cartesian Product is rarely what we want to produce...


Inner Joins

Types of Joins: two types


- *Inner joins* return only matching rows.




- *Outer joins* return all matching rows, plus nonmatching rows from one or both tables.



Left



Full



Right

Inner Join

Generate a report showing all valid order information:

| ID | Name | ID | Action | Amount |
|-----|-------|-----|----------|--------|
| 101 | Smith | 102 | Purchase | \$100 |
| 101 | Smith | 103 | Return | \$52 |
| 101 | Smith | 105 | Return | \$212 |
| 104 | Jones | 102 | Purchase | \$100 |
| 104 | Jones | 103 | Return | \$52 |
| 104 | Jones | 105 | Return | \$212 |
| 102 | Blank | 102 | Purchase | \$100 |
| 102 | Blank | 103 | Return | \$52 |
| 102 | Blank | 105 | Return | \$212 |

Inner Join

The inner join clause links two (or more) tables by a relationship between two columns.

```
select *  
from customers, transactions  
where customers.ID=transactions.ID;
```

```
SELECT object-item<, ...object-item>  
FROM table-name, ... table-name  
WHERE join condition  
         <AND sql-expression>  
         <other clauses>;
```

Abbreviating the Code

- A *table alias* is a temporary, alternative name for a table. You can make the query easier to read by using table aliases.

```
SELECT alias-1.object-item<, ...alias-2.object-item>  
  FROM table-name <AS> alias-1,  
        table-name <AS> alias-2  
  WHERE join-condition(s)  
  <other clauses>;
```

- The AS keyword is optional in the table alias syntax.

Abbreviating the Code

```
proc sql;  
select c.ID, Name, Action, Amount  
       from customers as c, transactions as t  
       where c.ID=t.ID;  
quit;
```



| ID | Name | Action | Amount |
|-----|-------|----------|--------|
| 102 | Blank | Purchase | \$100 |

Alternative Join Syntax

This alternative syntax names the join type and includes an ON clause

```
select c.ID, Name, Action, Amount
       from customers as c
       inner join
       transactions as t
       on c.ID=t.ID;
```

```
SELECT object-item <, ...object-item>
FROM table-name <<AS> alias>
INNER JOIN
       table-name <<AS> alias>
ON join-condition(s)
WHERE sql-expression
       <other clauses>;
```



Tables: **jupiter.employees**

jupiter.employee_addresses

Display: Employee_ID, Gender, Employee_Name



Tables: **practice.movies**


practice.genres

Display: movie_name, genre name


Outer Joins

Outer Joins


- *Inner joins* return only matching rows.




- *Outer joins*: you can retrieve both non-matching and matching rows using an outer join. Many tables can be referenced in outer joins. The tables are processed two tables at a time.



Left




Full



Right

Outer Joins

Generate a report that displays ***all*** customers and any transactions that they have completed.



| ID | Name | ID | Action | Amount |
|-----|-------|-----|----------|--------|
| 101 | Smith | . | | . |
| 102 | Blank | 102 | Purchase | \$100 |
| 104 | Jones | . | | . |

Outer Joins

Outer join syntax is similar to the alternate inner join syntax.

```
select *  
  from customers as c  
    left join  
    transactions as t  
  on c.ID=t.ID;
```

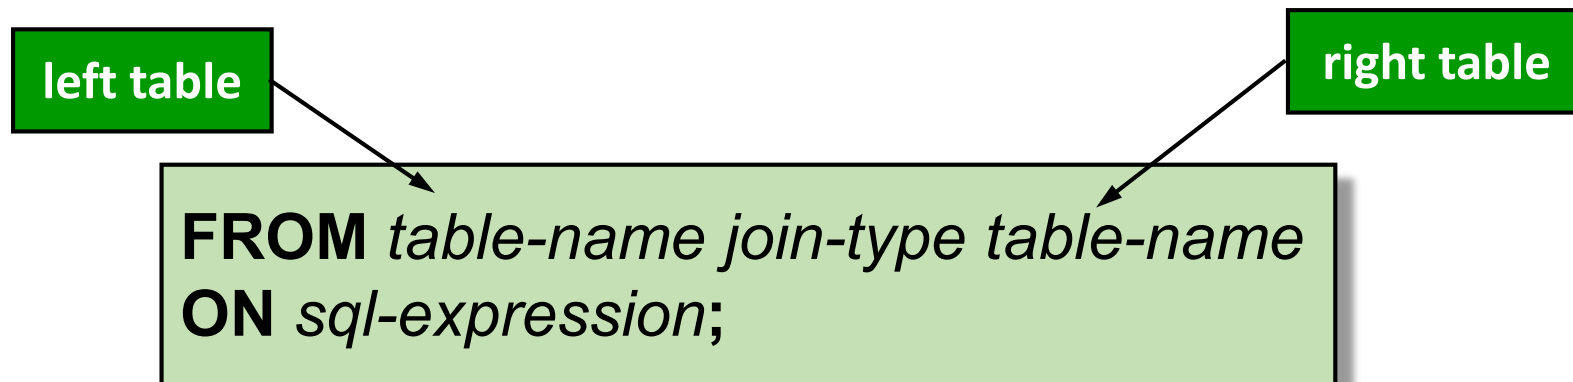
The ON clause
specifies the join
criteria in outer
joins.

```
SELECT object-item <, ...object-item>  
FROM table-name <<AS> alias>  
      LEFT|RIGHT|FULL JOIN  
      table-name <<AS> alias>  
ON join-condition(s)  
    <other clauses>;
```

Determining Left and Right

Consider the position of the tables in the FROM clause.

- Left joins return all matching and non-matching rows from the *left* table and the matching rows from the *right* table.
- Right joins return all matching and non-matching rows from the *right* table and the matching rows from the *left* table.
- Full joins return all matching and non-matching rows from all of the tables.



Left Join

customers

| ID | Name |
|-----|-------|
| 101 | Smith |
| 104 | Jones |
| 102 | Blank |

transactions

| ID | Action | Amount |
|-----|----------|--------|
| 102 | Purchase | \$100 |
| 103 | Return | \$52 |
| 105 | Return | \$212 |

```
select *  
  from customers c left join transactions t  
    on c.ID = t.ID;
```

| ID | Name | ID | Action | Amount |
|-----|-------|-----|----------|--------|
| 101 | Smith | . | . | . |
| 102 | Blank | 102 | Purchase | \$100 |
| 104 | Jones | . | . | . |

Includes all rows from the left table, even if there are no matching rows in the right table.

Right Join

customers

| ID | Name |
|-----|-------|
| 101 | Smith |
| 104 | Jones |
| 102 | Blank |

transactions

| ID | Action | Amount |
|-----|----------|--------|
| 102 | Purchase | \$100 |
| 103 | Return | \$52 |
| 105 | Return | \$212 |

```
select *  
  from customers c right join transactions t  
    on c.ID = t.ID;
```

| ID | Name | ID | Action | Amount |
|-----|-------|-----|----------|--------|
| 102 | Blank | 102 | Purchase | \$100 |
| . | | 103 | Return | \$52 |
| . | | 105 | Return | \$212 |

Includes all rows from the right table, even if there are no matching rows in the left table.

Full Join

customers

| ID | Name |
|-----|-------|
| 101 | Smith |
| 104 | Jones |
| 102 | Blank |

transactions

| ID | Action | Amount |
|-----|----------|--------|
| 102 | Purchase | \$100 |
| 103 | Return | \$52 |
| 105 | Return | \$212 |

```
select *  
  from customers c full join transactions t  
    on c.ID = t.ID;
```

| ID | Name | ID | Action | Amount |
|-----|-------|-----|----------|--------|
| 101 | Smith | . | . | . |
| 102 | Blank | 102 | Purchase | \$100 |
| . | | 103 | Return | \$52 |
| 104 | Jones | . | . | . |
| . | | 105 | Return | \$212 |

Includes all rows
from both tables,
even if there are no
matching rows in
either table

For the lab...

